

Yandex



Effective Structured Data Marshalling/Demarshalling Through Boost.Fusion Introspection In A High Performance Web Service

Piotr Reznikov, Sergei Khandrikov

The Effective Structured Data Marshalling/Demarshalling Through Boost.Fusion Introspection In A High Performance



Who We Are



Hello, we are Yandex

Yandex is the leading search engine in Russia.

- › 62% of Russian search traffic
- › 25 mln unique users per day
- › 6,000+ employees

We do mail, as well

- | Yandex.Mail is a free mail service, quite popular in Russia and Russian-speaking countries.
- › Built in 2000
- › 9 million unique users per day
- › 110 million messages sent and received daily

The Effective Structured Data Marshalling/Demarshalling Through Boost.Fusion Introspection In A High Performance



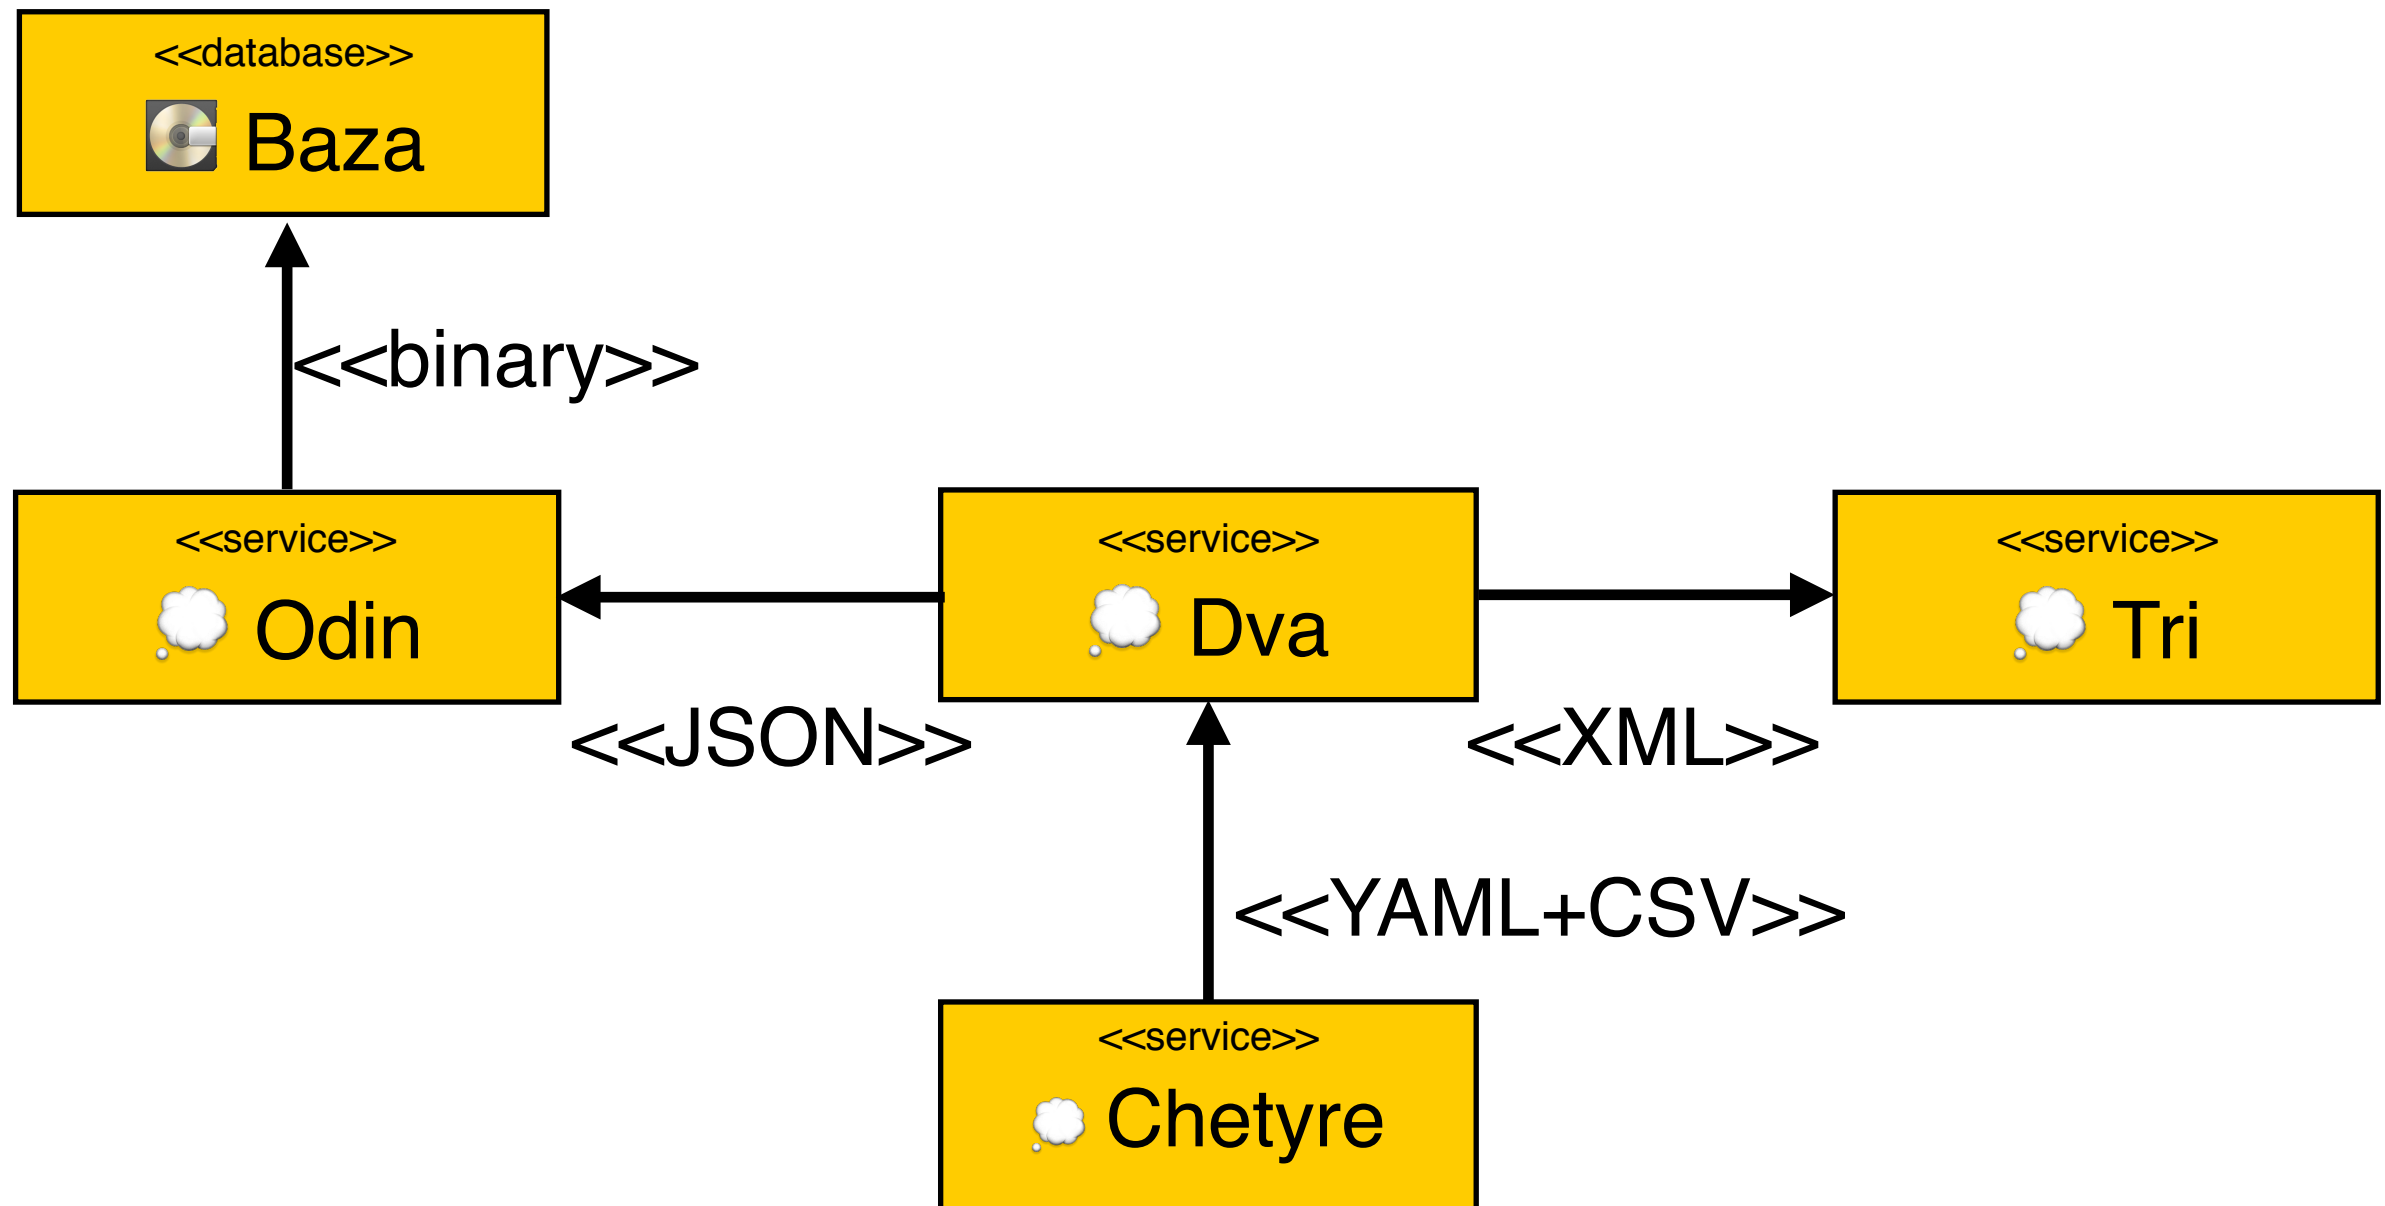
Problem



Situation

- | There are dozens of services inside and outside Mail which are:
 - › written in C++
 - › communicating via HTTP protocol with different text formats
 - › getting a data in binary format from a database

Situation



Situation

- | So, for various formats, we want
 - › Unification - convenient mechanism for C++.
 - › Optimisation - maximum performance.

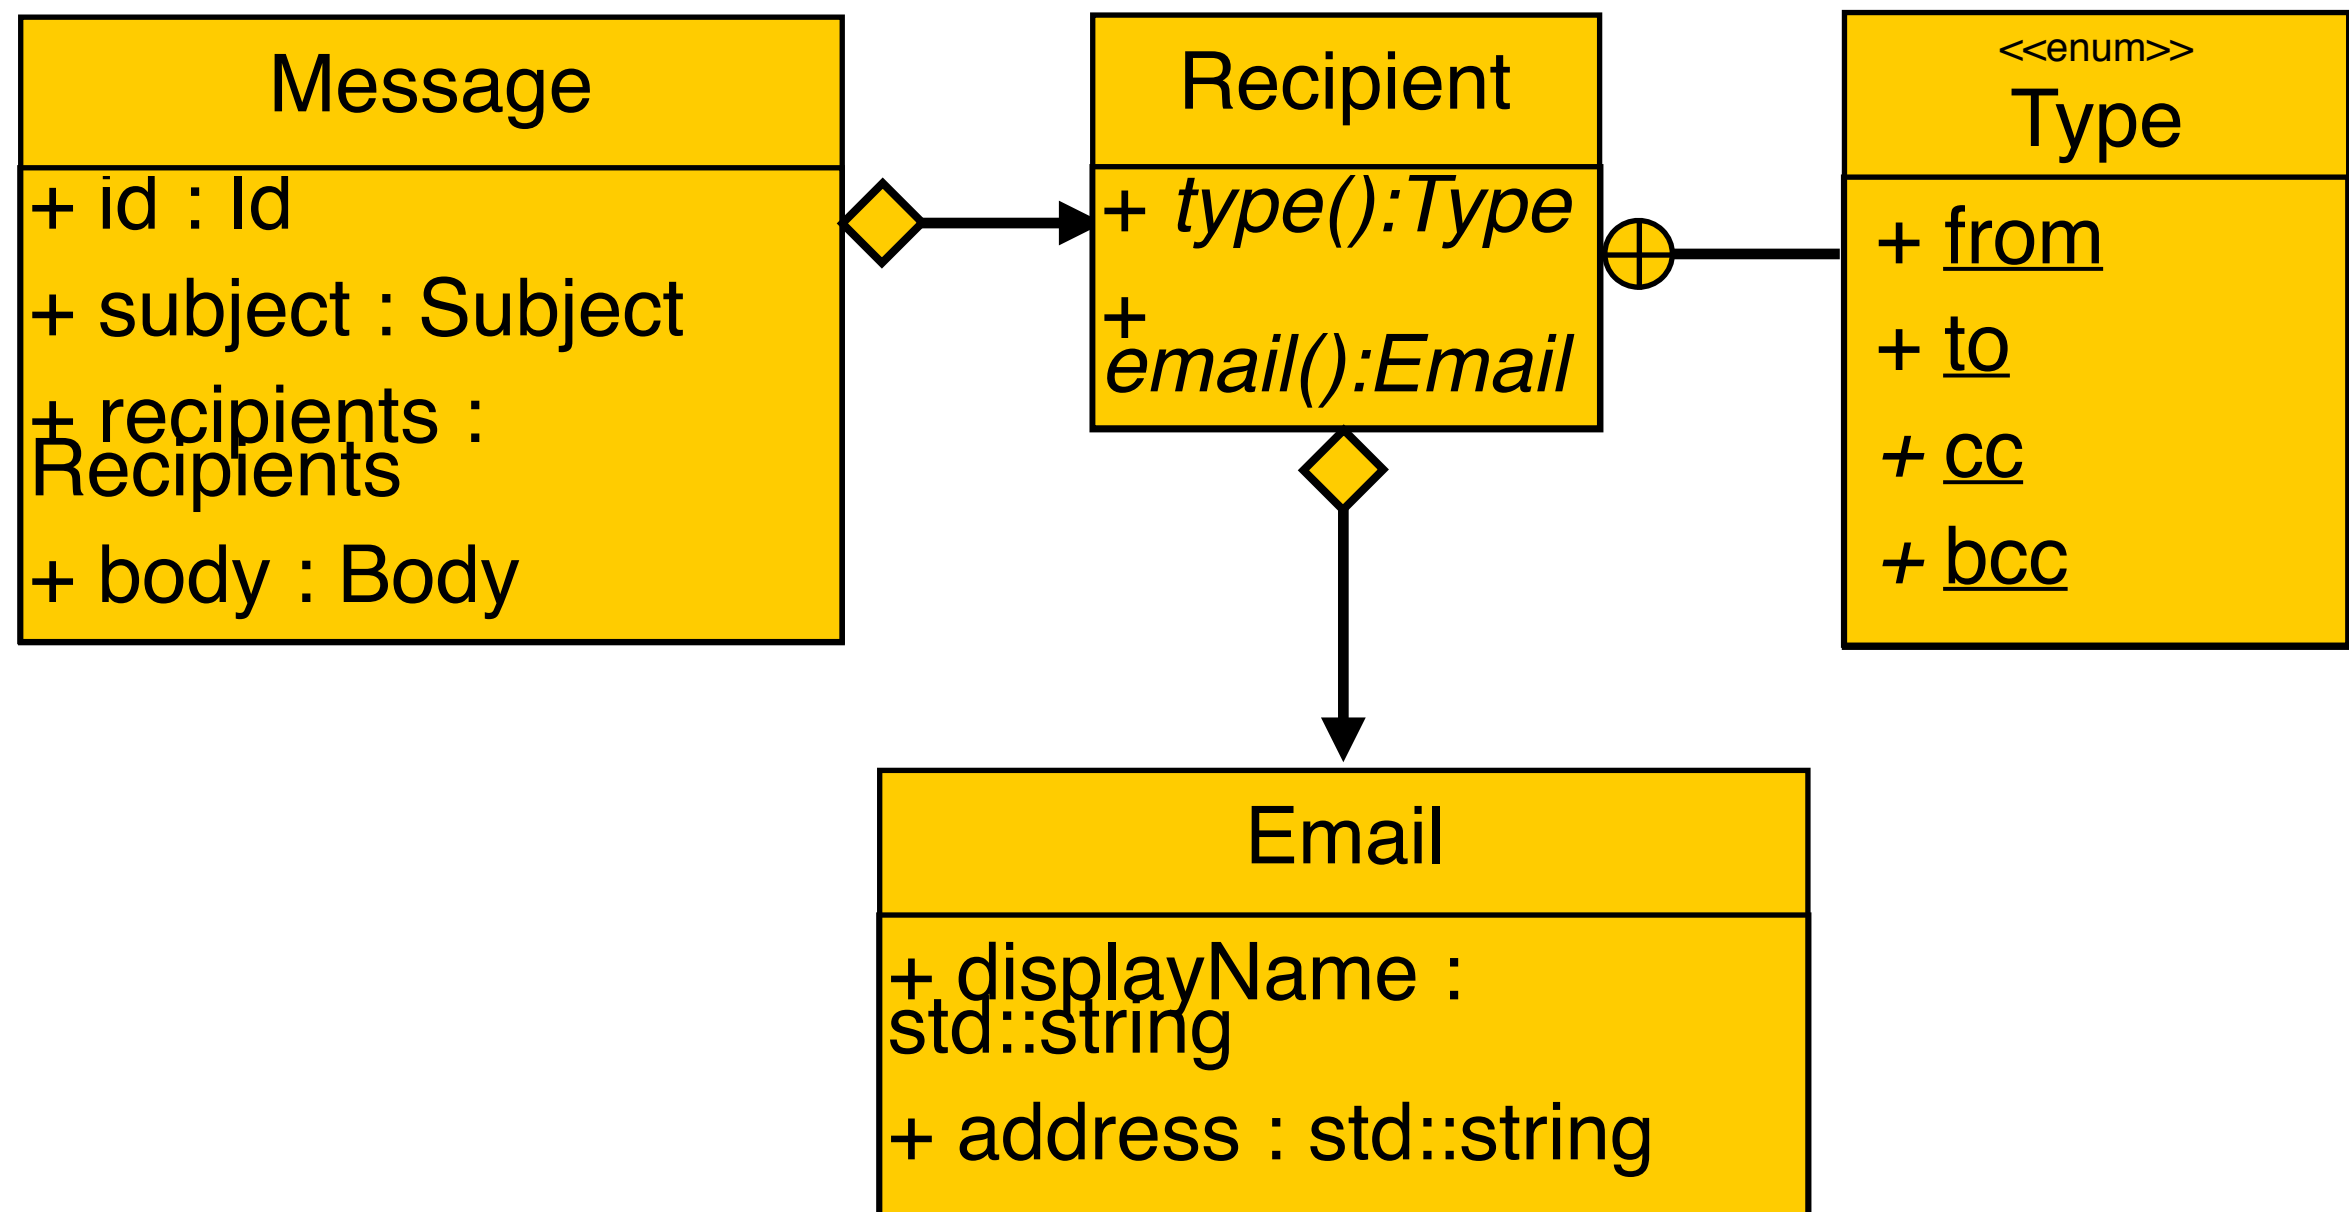
We need to serialize some struct

| Here's some trivialized e-mail class system:

```
struct Message {  
    std::string id;  
    std::string subject;  
    std::string body;  
    std::vector<Recipient> recipients;  
};
```

```
struct Recipient {  
    enum class Type { ... };  
    Type type() const;  
    Email email() const;  
};  
  
struct Email {  
    std::string name;  
    std::string address;  
};
```

We need to serialize some struct



Serialize into JSON

| So let's just put some YAJL here, right?

Serialize into JSON

So let's just put some YAJL here, right?

```
#define YAJL_ADD_FIELD_NAME(gen, name) \
    const unsigned char * const name##Name = reinterpret_cast<const unsigned char*>(#name); \
    yajl_gen_string(gen, name##Name, strlen(#name));

#define YAJL_ADD_STRING(gen, str) \
    yajl_gen_string(gen, reinterpret_cast<const unsigned char*>(str.c_str()), str.size());

#define YAJL_ADD_FIELD(gen, obj, field) \
    YAJL_ADD_FIELD_NAME(gen, field) \
    YAJL_ADD_STRING(gen, obj.field)

void print(yajl_gen_t* gen, const model::Email& e) {
    yajl_gen_map_open(gen);
    YAJL_ADD_FIELD(gen, e, name);
    YAJL_ADD_FIELD(gen, e, address);
    yajl_gen_map_close(gen);
}

void print(yajl_gen_t* gen, const model::Recipient& r) {
    yajl_gen_map_open(gen);
    YAJL_ADD_FIELD_NAME(gen, type);
    YAJL_ADD_STRING(gen, r.type());
    YAJL_ADD_FIELD_NAME(gen, email);
    print(gen, r.email());
    yajl_gen_map_close(gen);
}

void print(yajl_gen_t* gen, const model::Message& m) {
    yajl_gen_map_open(gen);

    YAJL_ADD_FIELD(gen, m, id);
    YAJL_ADD_FIELD(gen, m, body);
    YAJL_ADD_FIELD(gen, m, subject);

    YAJL_ADD_FIELD_NAME(gen, recipients);

    yajl_gen_array_open(gen);
    for( const auto& r : m.recipients ) {
        print(gen, r);
    }
    yajl_gen_array_close(gen);

    yajl_gen_map_close(gen);
}
```

```
std::string toJson(Message m) {
    yajl_gen_t* gen = yajl_gen_alloc(nullptr);
    yajl_gen_map_open(gen);
    YAJL_ADD_FIELD(gen, m, id); //we try our best to make that
code look nice
    YAJL_ADD_FIELD(gen, m, body);
    YAJL_ADD_FIELD(gen, m, subject);
    YAJL_ADD_FIELD_NAME(gen, recipients);
    yajl_gen_array_open(gen);
    for( const auto& r : m.recipients ) {
        //...
    }
    yajl_gen_array_close(gen);
    yajl_gen_map_close(gen);
}
```

Serialize into JSON and XML

- | But hey, you are only need to write this once. Or are you, really?
How about XML?

Deserialize from database format

| Another hundreds lines of code?

Serialize without «body» field

- | Exactly like before, but without heavy «body» field

Serialize without «body» field

| Which solution do you prefer?

› `std::string toJson(Message m, bool useBody);` *//non-extensible*

› `std::string toJson(Message m);` *//fat copy-paste*
`std::string toJsonNoBody(Message m);`

› `std::string toJson(Message m);` *//DTO, copy-paste*
`std::string toJson(MessageNoBody m);`

| We like neither.

Another caveats in handwritten solution

- › Repeat for every struct you want to serialize
- › Inverse-repeat for deserialize (for every struct, for every format)

We want pretty solution like this

```
JsonWriter jWriter;
```

```
Message msg;
```

```
auto json = jWriter.apply(msg);
```

```
template<typename T>  
string JsonWriter::apply(T t);
```

We want pretty solution like this

```
DBReader dbReader;    //deserializator for binary DB protocol
JsonWriter jWriter;

auto data = db.getMessage();

auto msg = dbReader.apply<Message>(data); //parse binary format
msg.adjustRecipients();                  //make adjustments to object
auto json = jWriter.apply(msg);          //give it away as json
```

We want pretty solution like this

```
JsonWriter jWriter;
```

```
Message msg;
```

```
auto json = jWriter.apply(msg);
```

```
struct AnOther { int x; };
```

```
ADAPT_STRUCT(AnOther, (int, x)); //Sad but true: no native  
                                //introspection in C++
```

```
auto jsonRec = jWriter.apply(AnOther()); //Same mechanism for any type
```

We want pretty solution like this

```
JsonWriter jWriter;
```

```
Message msg;
```

```
auto json = jWriter.apply(msg);
```

```
ADAPT_VIEW(Message, MessageNoBodyView, ...)
```

```
auto jsonNoBody = jWriter.apply<MessageNoBodyView>(msg); //Same type,  
//different view. No copies made
```

The Effective Structured Data Marshalling/Demarshalling
Through Boost.Fusion Introspection In A High Performance



The YReflection



Research subjects for the new solution

- › Easy definition - no external tools
- › Native - no Data Transfer Objects
- › Documented
- › Zero copy support

Basic Idea

- | Inspired by the Stack Overflow answer from Seth Heeren (SeHe) [\[1\]](#). Walking through the objects whose types are defined with Boost.Fusion
- › Define meta information with Boost.Fusion
- › Visiting data entities
- › Apply appropriate visitor for each attribute or method

We want pretty solution like this

```
JsonWriter jWriter;
```

```
Message msg;
```

```
auto json = jWriter.apply(msg);
```

Define Metadata With Boost.Fusion

| So it might look like this:

```
struct Message {  
    std::string id;  
    std::string subject;  
    std::string body;  
    RecipientList recipients;  
};  
  
BOOST_FUSION_ADAPT_STRUCT(Message,  
    id,  
    subject,  
    body,  
    recipients  
)
```

```
BOOST_FUSION_ADAPT_ADT(  
    Recipient,  
    (obj.type(), obj.type(val))  
    (obj.email(), obj.email(val))  
)  
  
BOOST_FUSION_ADAPT_STRUCT(  
    Email,  
    name,  
    address  
)
```

Define Metadata With Boost.Fusion

- | We want to define metadata using the Boost.Fusion Adapted like
 - › BOOST_FUSION_ADAPT_STRUCT //For existing type
 - › BOOST_FUSION_ADAPT_ADT //For type with setters/getters
 - › BOOST_FUSION_DEFINE_STRUCT //To define brand new struct

| New struct may be defined without double type of members:

```
BOOST_FUSION_DEFINE_STRUCT(  
    Email,  
    (std::string, name)  
    (std::string, address)  
)
```

What if the data type being changed

| We want to use BOOST_FUSION_ADAPT_NAMED:

```
//Lightened version of response
BOOST_FUSION_ADAPT_STRUCT_NAMED( Message, MessageNoBodyView,
    id, subject, recipients
)
```

| JSON:

```
{
  "id": "42-100500",
  "subject": "I love you Ozzy!",
  "recipients": [ ... ]
}
```

The documentation

| We want to add inline documentation looks like this:

```
YREFLECTION_ADAPT_DOC( Message, 'User mail message.',  
    (id, 'An unique message id.')  
    (subject, 'Subject header of a message.')  
    (body, 'Message text.')  
    (recipients, 'Message recipients.') )
```

| So we want to get online documentation looks like this:

```
$ curl "http://service/messages/help?format=yaml"  
Message          # User mail message.  
id: int          # An unique message id.  
subject: Subject # Subject header of a message.  
body: string     # Message text.  
recipients: vector<Recipient> # Message recipients.
```

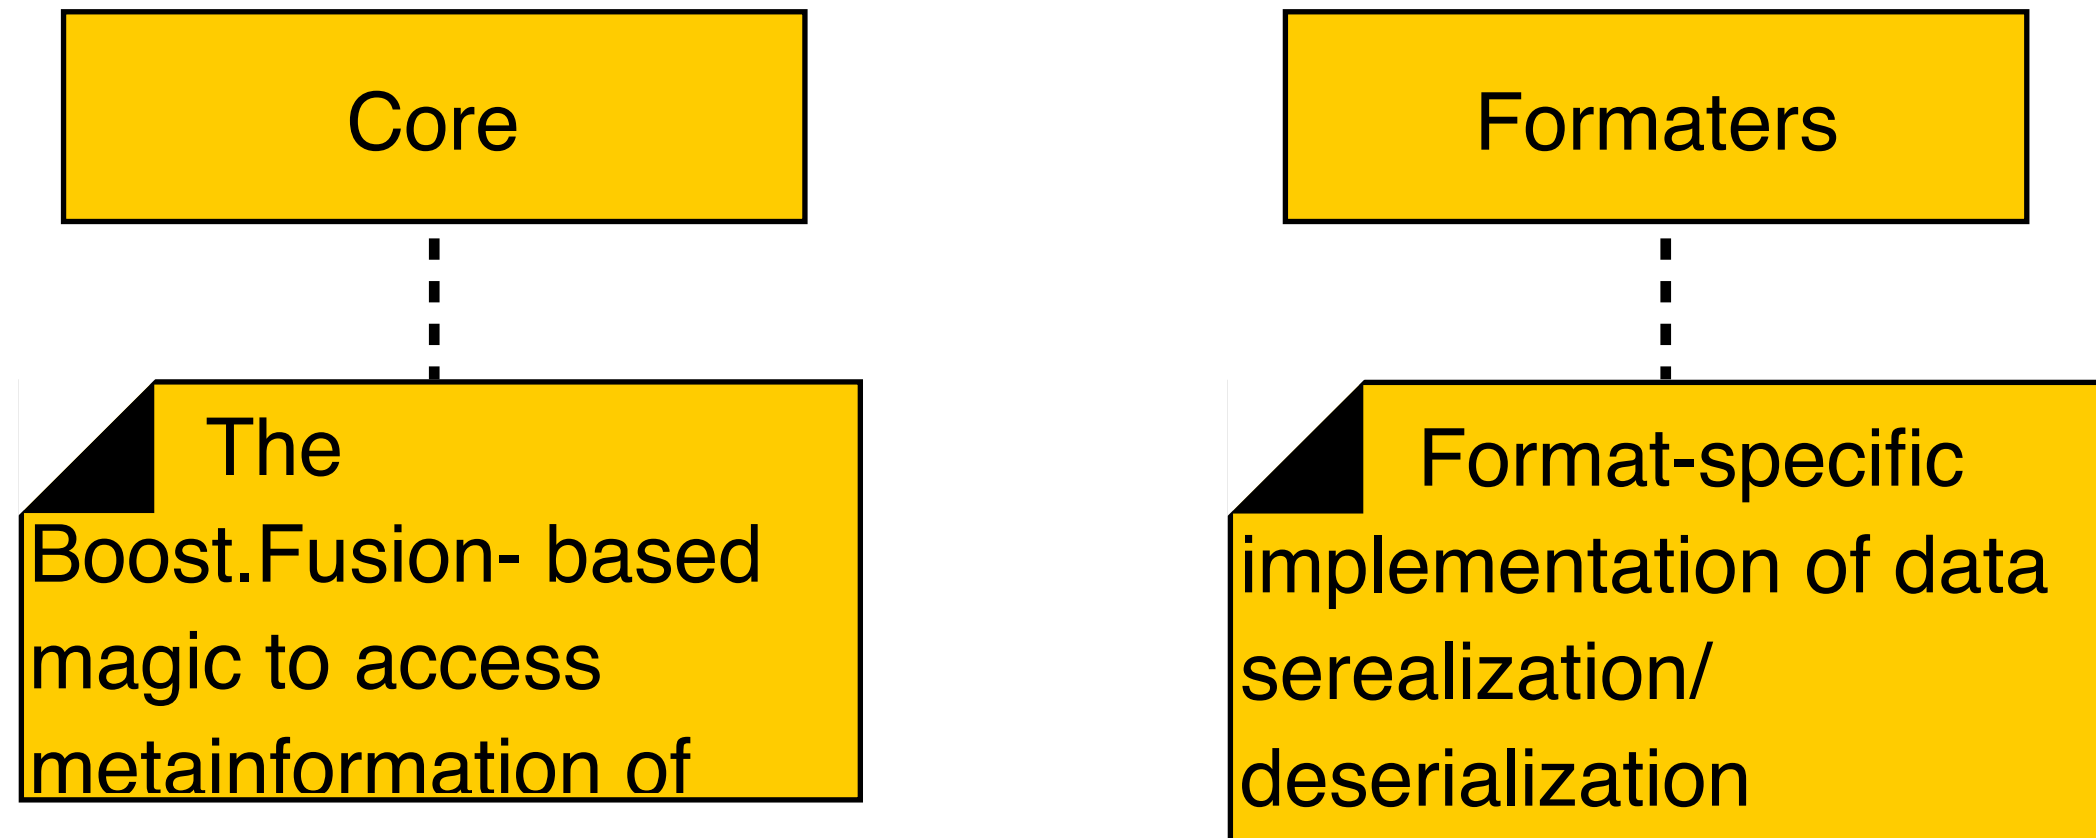
The Effective Structured Data Marshalling/Demarshalling Through Boost.Fusion Introspection In A High Performance



Architecture



The Architecture



The Core Magic Example

| When you do this:

```
Message msg;  
auto json = JsonWriter().apply(msg);
```

Core Magic Example

| When you do this:

Message msg;
`auto` json = JsonWriter().apply(msg);

```
string JsonWriter::apply(T t) {  
    core::apply(t, *this);  
    return yajl.result();  
}
```

Core Magic Example

| When you do this:


```
Message msg;  
auto json = JsonWriter().apply(msg);
```

```
string JsonWriter::apply(T t) {  
    core::apply(t, *this);  
    return yajl.result();  
}
```


| The core meets Message structure and uses this code:

```
template <typename Tag>  
static void apply (T& field, Visitor& v, Tag tag) {  
    auto members = members::make_vector(field);  
    auto itemVisitor = v.onStructStart(field, tag);  
    boost::fusion::for_each(members, visit_struct::adapt(itemVisitor));  
    v.onStructEnd(field, tag);  
}
```

Architecture



Visitor



JsonVisitor

- + onField(val : Field&&, tag : Tag)
- + onStructStart(s : Struct&&, tag : Tag) : Visitor
- + onStructEnd(s : Struct&&, tag : Tag)
- + onMapStart(s : Map&&, tag : Tag) : Visitor
- + onMapEnd(s : Map&&, tag : Tag)
- + onSequenceStart(s : Sequence&&, tag : Tag) :
Visitor
- + onSequenceEnd(s : Sequence&&, tag : Tag)
- + onOptional(v : Optional&&, tag : Tag) : bool

Visitor

| Field handling methods

```
template<typename V, typename Tag> void onField(V&& , Tag);
```

```
template<typename V, typename Tag> Visitor onStructStart(V&& , Tag);
```

```
template<typename V, typename Tag> void onStructEnd(V&& , Tag);
```

```
template<typename V, typename Tag> Visitor onMapStart(V&& , Tag);
```

```
template<typename V, typename Tag> void onMapEnd(V&& , Tag);
```

```
template<typename V, typename Tag> Visitor onSequenceStart(V&& , Tag);
```

```
template<typename V, typename Tag> void onSequenceEnd(V&& , Tag);
```

```
template<typename V, typename Tag> bool onOptional(V&& , Tag);
```

```
template<typename V, typename Tag> bool onSmartPointer(V&& , Tag);
```

Visitor

| Field's tags

```
struct MapItemTag;
```

```
struct SequenceItemTag;
```

```
template <typename Name>  
struct NamedItemTag;
```

```
template <typename ... Args>  
auto name(const NamedItemTag<Args...>& tag);
```

Visitor

| Plain Field method

```
template<typename V, typename Tag> void onField(V&& , Tag);
```

| Actually a group of specialized method

```
template<typename V, typename ... Arg>
void onField(const V& v, NamedItemTag<Args...> tag) {
    addJsonMapItem(name(tag), boost::lexical_cast<std::string>(v));
}
```

```
template<typename V>
void onField(const V& v, SequenceItemTag) {
    addJsonArrayItem(boost::lexical_cast<std::string>(v));
}
```


The Visitor Concept

| Classes handle methods

```
template<typename V, typename Tag> Visitor onStructStart(V&& , Tag);  
template<typename V, typename Tag> void onStructEnd(V&& , Tag);
```

```
template<typename V, typename Tag> Visitor onMapStart(V&& , Tag);  
template<typename V, typename Tag> void onMapEnd(V&& , Tag);
```

```
template<typename V, typename Tag> Visitor onSequenceStart(V&& , Tag);  
template<typename V, typename Tag> void onSequenceEnd(V&& , Tag);
```

The Visitor Concept

| Classes handle methods

```
template<typename V, typename ... Args>
Visitor onStructStart(const V& , NamedItemTag<Args...> tag) {
    openJsonMap(name(tag));
    return Visitor(jsonHandle);
}
```

```
template<typename V, typename Tag>
void onStructEnd(const V& , Tag) {
    closeJsonMap();
}
```

The Visitor Concept

| Return Visitor allows handle mixed format

```
template<typename V, typename ... Args>
CSVReader JsonReader::onSequenceStart(CSVSequence& ,
                                     NamedItemTag<Args...> tag) {
    auto data = node(name(tag)).data();
    return CSVReader(data);
}
```

```
template<typename V, typename Tag>
void JsonReader::onSequenceEnd(const V& , Tag) {
}
```

The Visitor Concept

| Optional field handle methods

```
template<typename V, typename Tag> bool onOptional(V&& , Tag);  
template<typename V, typename Tag> bool onSmartPointer(V&& , Tag);
```

| Trivial implementation for serialization

```
template<typename V, typename Tag> bool onOptional(V const& v, Tag) {  
    return v.is_initialized();  
}
```

The Visitor Concept

| Trivial implementation for deserialization

```
template<typename V, typename ... Args>
bool onOptional(V& v, NamedItemTag<Args...> tag) {
    if (currentName() == name(tag)) {
        v = V();
        return true;
    }
    return false;
}
```

The Effective Structured Data Marshalling/Demarshalling
Through Boost.Fusion Introspection In A High Performance



Competition



Competitors

- › *yajl* - hand-made serialization via yajl library and simple http server
- › *YReflection* - serialization via yreflection and simple http server
- › *Protobuf* - http server with protobuf support, as a different approach

What do we measure

- › Average latency
- › CPU consumption
- › Memory consumption
- › Lines of code (LOC) as a code complexity

Hardware

- › Intel® Xeon® Processor E5645 (12M Cache, 2.40 GHz)
- › Memory 92 Gb
- › Net 2x 1000baseT-FD.

Software

- › Ubuntu 14.04.4 x86_64
- › gcc version 4.8.4
- › Boost 1.60
- › Fusion 2.2
- › yajl 2
- › Protobuf 2.5

The Effective Structured Data Marshalling/Demarshalling
Through Boost.Fusion Introspection In A High Performance

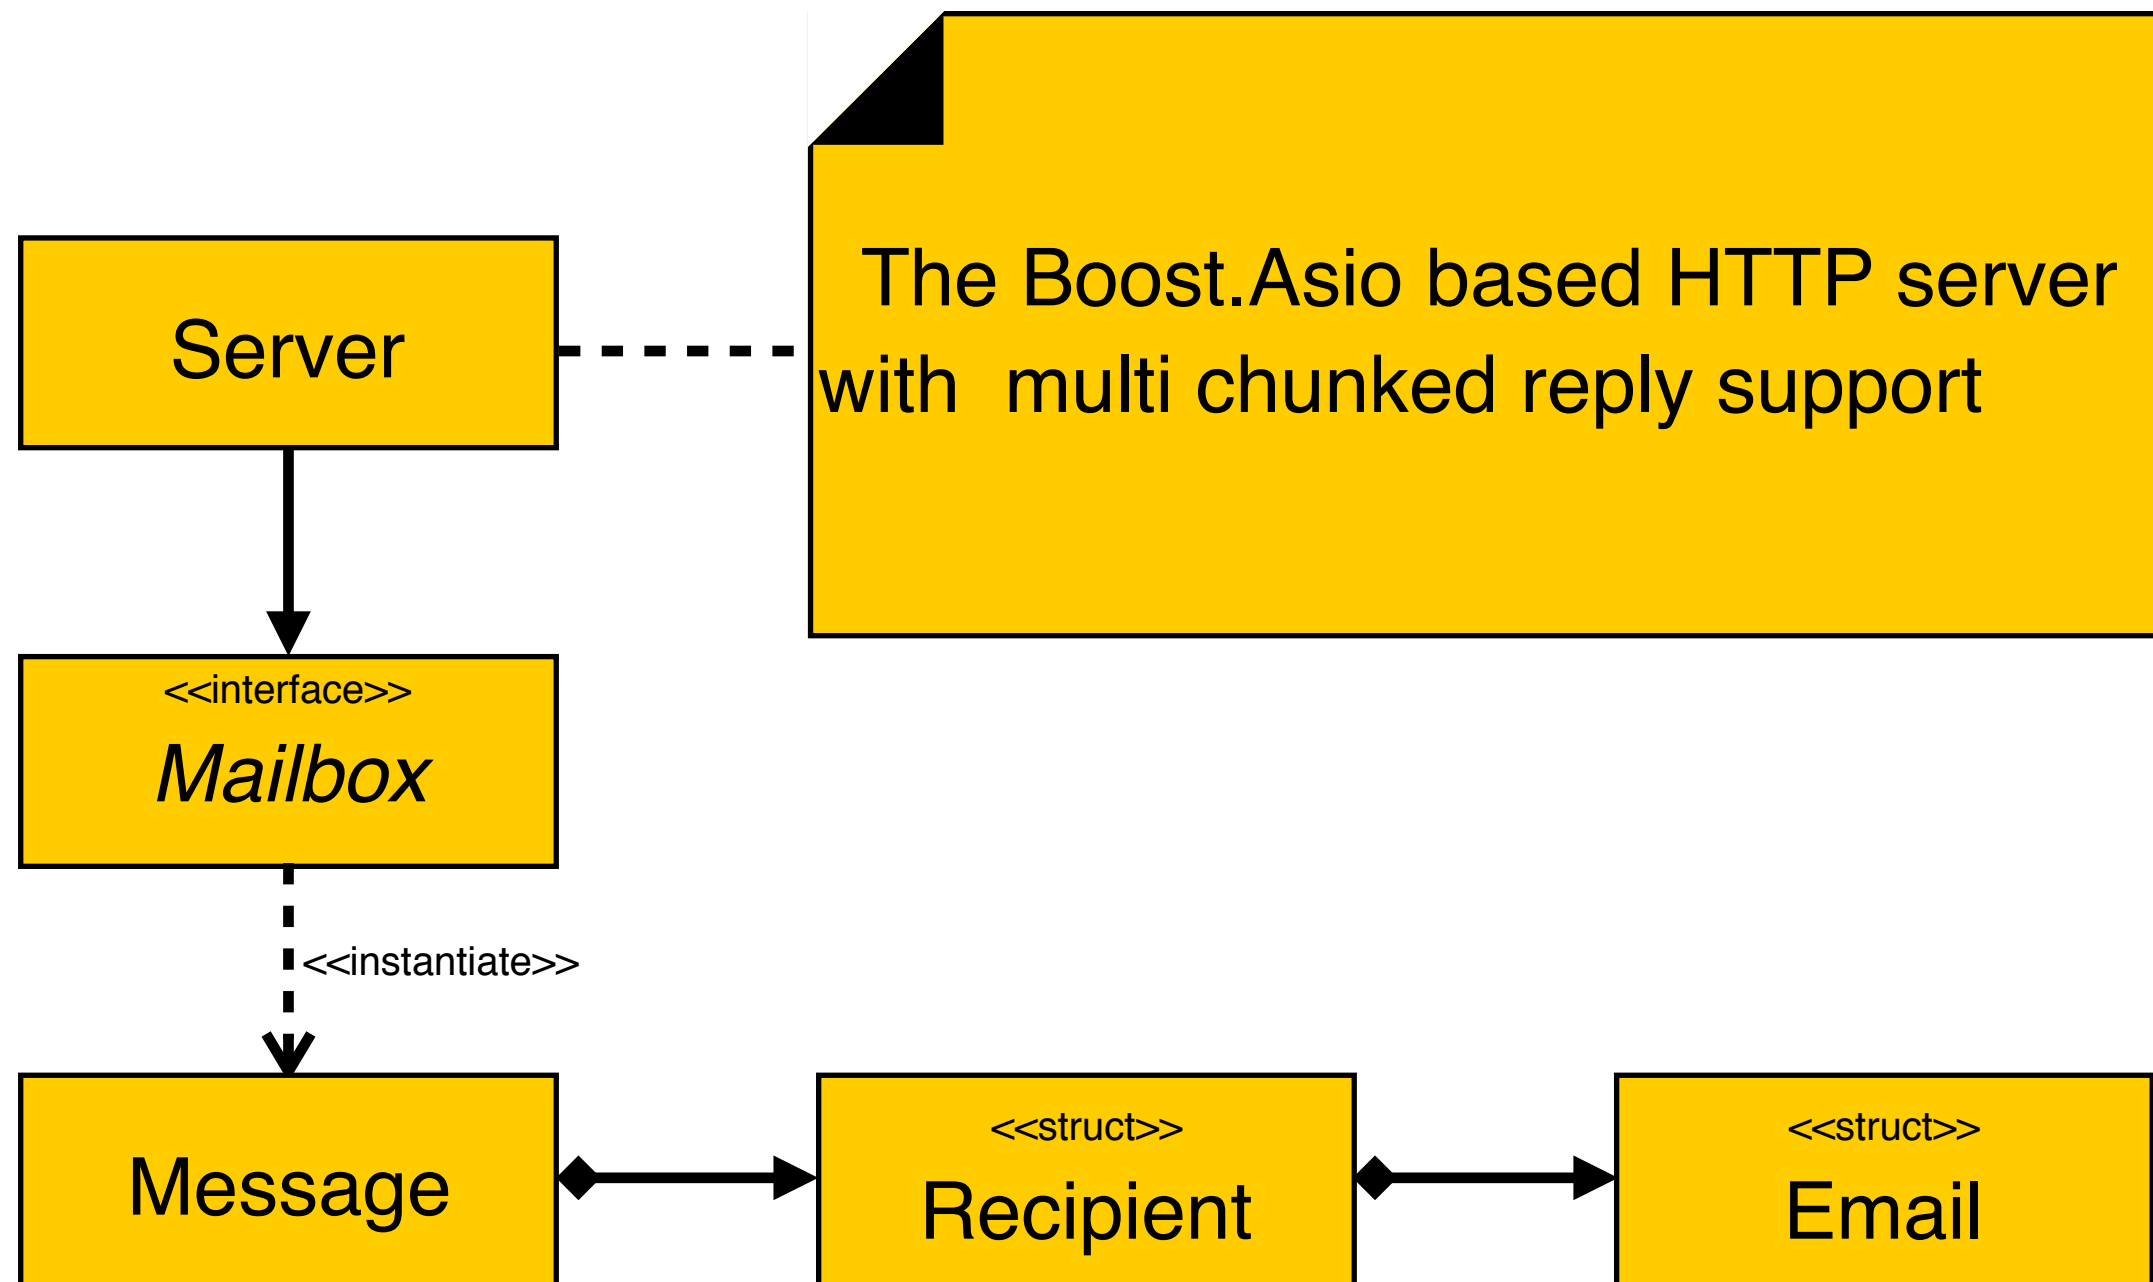
Model



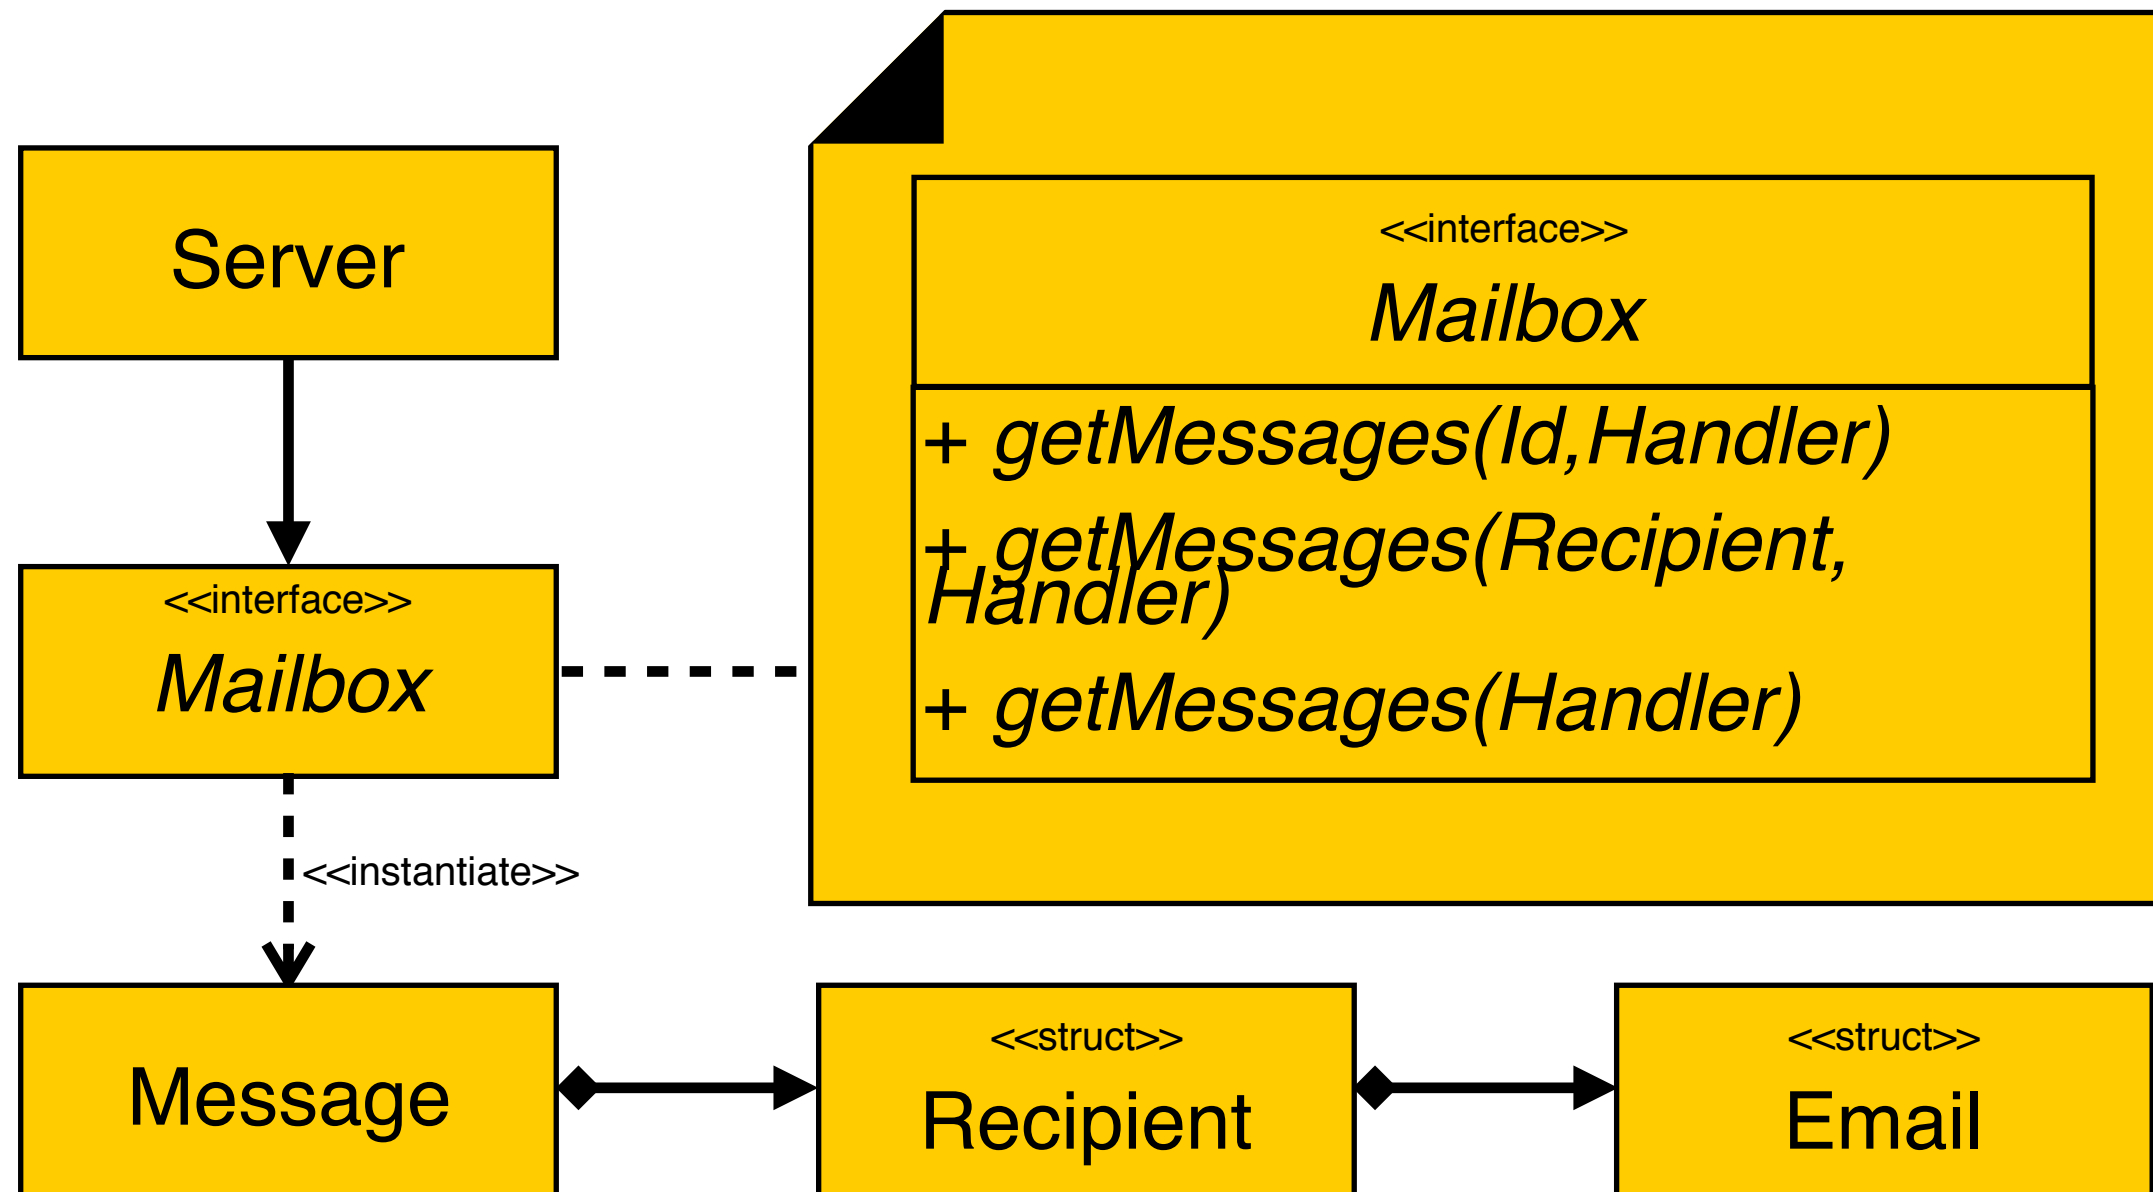
Mailbox access via web service model

- | We examine the web service which provides simple API to access user's mailbox
 - › *messages/* - return all the mailbox messages
 - › *messages/id* - returns a message by id
 - › *messages/recipient* - returns messages by recipient - **used in benchmark**


The Service Architecture



The Service Architecture




The Effective Structured Data Marshalling/Demarshalling Through Boost.Fusion Introspection In A High Performance




Results



Latency, ms



CPU consumption



Memory, KB



LOC

yreflection




250

yajl




70

protobuf



75

LOC



The Effective Structured Data Marshalling/Demarshalling
Through Boost.Fusion Introspection In A High Performance

The Boost.Fusion Issues



Defining different versions of response

| Let's try to use BOOST_FUSION_ADAPT_NAMED:

```
struct Struct {  
    Array array;  
    int id;  
};  
struct Response {  
    Struct b;  
};
```

```
BOOST_FUSION_ADAPT_STRUCT(Struct, id, array)  
BOOST_FUSION_ADAPT_STRUCT(Response, b)  
//Old version of response  
BOOST_FUSION_ADAPT_STRUCT_NAMED(Struct, StructV1, id)  
BOOST_FUSION_ADAPT_STRUCT_NAMED(Response, ResponseV1, b)
```

Something goes wrong

| Oops, we got the warning:

```
reflection/struct.cc:105:1: warning: returning reference to local  
temporary object  
[-Wreturn-stack-address]
```

access::struct_member::apply() problem

| Simplified macro expansion of BOOST_FUSION_ADAPT_NAMED

```
struct access::struct_member< ResponseV1 , 0 > {  
    typedef StructV1 attribute_type;  
    template<typename Seq>  
    struct apply {  
        typedef typename const_if<attribute_type, Seq> inner_type;  
        typedef typename add_reference<inner_type>::type type;  
        //yes, it adds reference always  
  
        constexpr static type call(Seq& seq) {  
            return seq.obj.b; //for a view we return a ref to the local  
        }  
    };  
};
```


The solution is

| We could return by-field when «Seq» type is a view, like this

```
struct access::struct_member< ResponseV1 , 0 > {
    typedef StructV1 attribute_type;
    template<typename Seq>
    struct apply {
        typedef typename const_if<attribute_type, Seq> inner_type;
        typedef typename mpl::eval_if< struct_is_view<attribute_type>,
            add_reference<inner_type>,
            inner_type
        >::type
    } type;
    constexpr static type call(Seq& seq) {
        return seq.obj.b; //for a view we return by field
    }
};
```

View problem

| Generated view cannot be used with const objects

```
BOOST_FUSION_ADAPT_STRUCT_NAMED( Struct, StructV1, id )
```

```
// ...
```

```
namespace {  
    struct StructV1 {  
        constexpr StructV1(Struct& in_obj) // No const support  
            : obj(in_obj) {}  
        Struct& obj;  
    private:  
        StructV1& operator= (StructV1 const&);  
    };  
}
```

Probable solution to View problem

| We can define additional view for const objects

```
BOOST_FUSION_ADAPT_STRUCT_NAMED( Struct, StructV1, id )  
BOOST_FUSION_ADAPT_STRUCT_NAMED( const Struct, StructCV1, id )  
// ...
```

```
namespace {  
    struct StructV1 {  
        constexpr StructV1(Struct& in_obj);  
    };  
  
    struct StructCV1 {  
        constexpr StructCV1(const Struct& in_obj);  
    };  
}
```

| Doesn't scale well when adapting nested structures

boost_fusion_adapt_adt_impl_set()

| Boost.Fusion doesn't support moving field into setter

```
class A {  
    void set(std::string val);  
};
```

```
BOOST_FUSION_ADAPT_ADT(A, (obj.get(), obj.set(forward<Val>(val))))
```

| Because part of macro expansion for setter looks like this:

```
template<class Val>  
static void boost_fusion_adapt_adt_impl_set(A& obj, Val const& val) {  
    obj.set(forward<Val>(val)); // No perfect forwarding, sorry  
}
```

Solution is

| Boost.Fusion doesn't support moving field into setter

```
class A {  
    void set(std::string val);  
};
```

```
BOOST_FUSION_ADAPT_ADT(A, (obj.get(), obj.set(forward<Val>(val))))
```

| Because part of macro expansion for setter looks like this:

```
template<class Val>  
static void boost_fusion_adapt_adt_impl_set(A& obj, Val && val) {  
    obj.set(forward<Val>(val)); // Perfect forwarding works here  
}
```

BOOST_FUSION_ADAPT_ADT - No

| You can't link a name to getter/setter pair

```
struct A {  
    void field(std::string const& val);  
    std::string const& field() const;  
};
```

```
BOOST_FUSION_ADAPT_ADT(A, (obj.field(), obj.field(val)))  
auto json = JsonWriter.apply(A("foo"), "A");
```

| What we want:

```
{  
  "A": {  
    "field": "foo"  
  }  
}
```

| What we get:

```
{  
  "A": [  
    "foo"  
  ]  
}
```

First non-optimal solution

- | We emulate named field by proxying it with `std::pair` of name and value

```
inline std::string stripMethodName(std::string name);

#define YR_CALL_WITH_SPECIFIC_NAME(fun, name) \
    std::make_pair(name, obj.fun())

#define YR_CALL_WITH_NAME(fun) \
    YR_CALL_WITH_SPECIFIC_NAME(fun, stripMethodName(#fun))

#define YR_CALL_SET_WITH_NAME(fun) obj.fun( val.second );

BOOST_FUSION_ADAPT_ADT(A,
    (YR_CALL_SET_WITH_NAME(field), YR_CALL_WITH_NAME(field)) )
```

- | Copy of name (and function result, if given by ref) made; fragile

A near-optimal solution is

| The solution by Sehe can be found on Stack Overflow [\[1\]](#)

```
#define MY_ADT_MEMBER_NAME(CLASSNAME, IDX, MEMBERNAME)\
    namespace boost { namespace fusion { namespace extension {\
    template <> struct struct_member_name<CLASSNAME, IDX> {\
        typedef char const* type;\
        static type call() { return #MEMBERNAME; }\
    }; }}}}
```

```
BOOST_FUSION_ADAPT_ADT(A, obj.field(), obj.field(val))
MY_ADT_MEMBER_NAME(A, 0, field)
```

| Have to declare it separately from main macro and point field's index in MPL sequence manually

What we want

| Automatic deduction from getter's name

```
struct A {  
    void field(std::string const& val);  
    std::string const& field() const;  
    std::string const& readOnlyField() const;  
};
```

```
YREFLECTION_ADAPT_ADT_PROPERTIES(A,  
    field, readOnlyField ) // Name can be deduced from getter name
```

| JSON:

```
{  
    "field": "foo"  
}
```

What we want

| Explicitly set field name

```
struct A {  
    void field(std::string const& val);  
    std::string const& field() const;  
    std::string const& readOnlyField() const;  
};
```

```
YREFLECTION_ADAPT_ADT_SET_GET(A,  
    (field, getField, setField)  
    (readOnlyField, getReadOnlyField) ) // Name is specified directly
```

| JSON:

```
{  
    "field": "foo",  
    "readOnlyField": "bar",  
}
```

What we want

| External set and get functions

```
struct A {  
    void field(std::string const& val);  
    std::string const& field() const;  
    std::string const& readOnlyField() const;  
};
```

```
Item str2item(std::string); // Converts string to the brand new Item  
std::string item2str(Item); // Converts Item to the old string
```

```
YREFLECTION_ADAPT_ADT_SET_GET(A,  
    (field, obj.setField(item2str(val)), str2item(obj.getField()))))
```

The Effective Structured Data Marshalling/Demarshalling Through Boost.Fusion Introspection In A High Performance



Summary



Conclusions

| Yreflection:

- › Has near-zero runtime overhead
- › Allows to write generic representation code
- › Is compact to use for clients

Summary

| We love Boost.Fusion

- › It gives us convenient mechanism for introspection
- › But it needs to be a little bit fixed and changed

| We want native introspection in C++

- › Boost.Fusion adaptation is good, but requires heavy macros
- › We don't want provide compile-time information to compiler

What we do next

- | Features to be implemented
 - › Documentation
 - › ADT with names
 - › Magic for SAX parsers

The Effective Structured Data Marshalling/Demarshalling
Through Boost.Fusion Introspection In A High Performance

 Questions?



Links

1. Stack Overflow post: <http://stackoverflow.com/questions/26380420/boost-fusion-sequence-type-and-name-identification-for-structs-and-class>
2. Boost.Fusion documentation: http://www.boost.org/doc/libs/1_60_0/libs/fusion/doc/html/index.html
3. Our dedicated project on Github: <https://github.com/thed636/reflection>

The Effective Structured Data Marshalling/Demarshalling
Through Boost.Fusion Introspection In A High Performance

 Спасибо (Thank
you)!



Contacts

Sergei Khandrikov

Senior Software Developer



thed636@yandex-team.ru



@S_Khandrikov



thed636

Contacts

Piotr Reznikov

Software Developer



prez@yandex-team.ru



@



pyotr.reznikov